




# MongoDB



## 1. Introduction to MongoDB

- What is MongoDB?
- NoSQL vs. SQL Databases: Understanding the Differences
- Key Features and Benefits of Using MongoDB

## 2. Getting Started with MongoDB

- Installing MongoDB
- Basic Configuration
- Understanding the MongoDB Environment

## 3. Fundamentals of MongoDB

- Databases and Collections
- Documents: The Basics
- CRUD Operations: Create, Read, Update, Delete

## 4. Working with Data in MongoDB

- Data Types and Schemas
- Inserting Data
- Querying Data: Basic and Advanced Queries
- Updating and Deleting Data

## 5. Indexes and Performance

- Understanding Indexes
- Creating and Managing Indexes
- Performance Tuning Tips

## 6. Data Aggregation

- Aggregation Framework: Concepts and Usage

- Pipeline Stages and Common Operators
- Practical Examples of Data Aggregation

7. Security in MongoDB

- Basic Security Practices
- Role-Based Access Control
- Encrypting Data

8. Scaling and Management

- Replication: Concepts and Configuration
- Sharding: Strategies and Implementation
- Monitoring and Maintenance

# 1. Introduction to MongoDB

---

- What is MongoDB?

MongoDB is a document-oriented database classified as a NoSQL database. Here's what that means:

- **Document-oriented:** Data in MongoDB is stored in flexible documents using a format similar to JSON called BSON (short for Binary JSON). Imagine these documents as enhanced versions of JSON objects where you can store arrays, other documents, and a variety of data types within.
- **NoSQL:** NoSQL stands for "Not Only SQL." These databases deviate from the traditional rigid, table-based structure of relational databases (more on this below).

- NoSQL vs. SQL Databases: Understanding the Differences

Feature	SQL (Relational Databases)	NoSQL (Document Databases)
Data Structure	Rigid tables with predefined columns and rows	Flexible documents (JSON-like) where fields can vary between documents
Schema	Schema-on-write: Enforced structure at the time of data entry	Schema-on-read: Structure is interpreted when data is read, allowing for more flexibility
Scaling	Typically scales vertically (adding more powerful hardware)	Scales horizontally well (adding more servers to a cluster)
Use Cases	Ideal for structured datasets with pre-defined relationships	Suited for rapidly changing data, semi-structured data, and applications needing

Feature	SQL (Relational Databases)	NoSQL (Document Databases)
		high scalability

- **Flexible Schema:** Documents don't need to adhere to a fixed structure. This is fantastic for evolving data models or handling data with some variability.
- **Scalability:** MongoDB excels at horizontal scaling. You can easily distribute data across multiple machines, making it well-suited for large-scale applications.
- **Performance:** Optimized for reads and writes, particularly when paired with effective indexing.
- **Rich Query Language:** You can perform complex queries, aggregations, and manipulate data within MongoDB itself.
- **Developer Friendly:** The JSON-like document model often aligns well with how developers represent data in their applications, easing interactions.

## 2. Getting Started with MongoDB

### Installation

There are two main paths depending on your goals:

- **MongoDB Atlas (Cloud-based):**
  - Ideal for quickly getting started and avoiding local setup.
  - Go to <https://www.mongodb.com/cloud/atlas> and create a free Atlas cluster. This handles installation and configuration behind the scenes.
- **MongoDB Community Edition (Local Installation):**
  - Great for development or learning on your own machine.
  - **Steps:**
    - Download:** Find the appropriate installer for your operating system on the MongoDB download page: <https://www.mongodb.com/try/download/community>
    - Run the installer:** Follow instructions for your OS.
    - Data Directory (Important):** MongoDB stores data in a default location (usually something like `/data/db`). Take note of this.

### Basic Configuration

- **Starting the MongoDB server:**
  - **Atlas:** Managed for you, no action needed!
  - **Local Install:**
    - Open a terminal / command prompt.
    - Navigate to the MongoDB installation directory (where `mongod` is located).
    - Run the command `mongod` (you might need to add options to specify the data directory you noted earlier).

## Understanding the MongoDB Environment

- **The Shell (mongosh):** This is your primary interaction point. Think of it as your command line for MongoDB.
  - **Atlas:** Provided within the web interface
  - **Local Install:** In a separate terminal, run the command `mongosh`.
- **Databases:** MongoDB organizes data into databases. Here's some common terminology:
  - **database** : A container for collections (similar to how a traditional database contains tables).
  - **collection** : A group of documents (analogous to a table in a relational database) .
  - **document** : A single record within a collection (like a row in a table).

## Essential Shell Commands (try these out in `mongosh` )

- `show dbs` : List existing databases.
- `use <database_name>` : Switch to a database ( `use mydatabase` ). You can create one if it doesn't exist.
- `db.mycollection.insertOne({ field: "value" })` : Insert a simple document into a collection named 'mycollection'.
- `db.mycollection.find()` : Retrieve all documents from the collection.

## Prerequisites

- **Have MongoDB set up:** Either a local installation or (easier for starting out) a MongoDB Atlas cluster.
- **Install PyMongo:** This is the official driver for interacting with MongoDB from Python. From your terminal run:

```
pip install pymongo
```

## Connecting to MongoDB

```
import pymongo

# For Atlas: Replace with your connection string
client = pymongo.MongoClient("mongodb+srv://<username>:<password>@<your-atlas-cluster>.mongodb

# Select database
db = client["mydatabase"]

# Select collection
collection = db["mycollection"]
```

## Advanced MongoDB Shell Commands in Python

- Insert:

```
result = collection.insert_one({ "name": "David", "interests": ["reading", "sports"] })
print(result.inserted_id) # Get the ID of the inserted document
```

- Find (Queries):

```
cursor = collection.find({"interests": "coding"}) # Find documents with 'coding' interest
for document in cursor:
    print(document)
```

- Update:

```
collection.update_one({"name": "David"}, {"$set": {"city": "New York"}})
```

- Delete:

```
collection.delete_many({"interests": {"$in": ["travel"]}})
```

- Aggregation:

```
pipeline = [
    {"$match": {"interests": "hiking"}},
    {"$group": {"_id": "$city", "totalHikers": {"$sum": 1}}}
]
results = collection.aggregate(pipeline)
for result in results:
    print(result)
```

## 3. Fundamentals of MongoDB in Python

### Prerequisites (Recap)

- MongoDB: Running locally or on Atlas.
- PyMongo: Installed ( `pip install pymongo` ).

### Databases and Collections

```
import pymongo

client = pymongo.MongoClient("...") # Your connection string
```

```
db = client["mydatabase"]

# Collections are implicitly created if they don't exist
mycollection = db["customers"]
```

- **Analogy:**
  - Database: A large container for related data.
  - Collection: A grouping of similar documents, similar to a table in SQL databases.

## Documents: The Basics

```
customer_document = {
    "name": "Emily",
    "address": {
        "street": "234 Oak Lane",
        "city": "Los Angeles"
    },
    "orders": [12345, 67890]
}
```

- **Key Points:**
  - Documents are like JSON objects (Python dictionaries).
  - Fields can have various data types (strings, numbers, booleans, arrays, subdocuments).
  - **\_id field:** A unique identifier automatically generated for each document.

## CRUD Operations: Create, Read, Update, Delete

- **Create (Insert):**

```
result = mycollection.insert_one(customer_document)
print(result.inserted_id)
```

- **Read (Find):**

```
# Find all:
cursor = mycollection.find()

# Find with a filter:
customer = mycollection.find_one({"name": "Emily"})

# Iterate over results:
for doc in cursor:
    print(doc)
```

- **Update:**

```

mycollection.update_one(
    {"name": "Emily"},
    {"$set": {"address.city": "New York"}}
)

# Update multiple documents
mycollection.update_many(
    {"orders": {"$exists": True}}, # Filter for those with an orders field
    {"$addToSet": {"orders": 99999}} # Add to the orders array
)

```

- Delete:

```

mycollection.delete_one({"name": "Emily"})

mycollection.delete_many({}) # Be careful! Deletes *all* documents

```

## Scenario: Managing a Book Collection

### Document Structure

```

book_document = {
    "title": "The Hitchhiker's Guide to the Galaxy",
    "author": "Douglas Adams",
    "genre": "Science Fiction",
    "published_year": 1979,
    "in_stock": True
}

```

### Python Script

```

import pymongo

# Connect to MongoDB (replace with your connection string)
client = pymongo.MongoClient("mongodb://localhost:27017/")
db = client["my_library"]
books_collection = db["books"]

# 1. Insert some documents
books = [
    {
        "title": "Pride and Prejudice",
        "author": "Jane Austen",
        "genre": "Classic",
        "published_year": 1813,
        "in_stock": True
    },

```

```

{
    "title": "To Kill a Mockingbird",
    "author": "Harper Lee",
    "genre": "Classic",
    "published_year": 1960,
    "in_stock": False
},
# ... add more if you'd like
]
books_collection.insert_many(books)

# 2. Query with different filters
print("Science Fiction Books:")
for book in books_collection.find({"genre": "Science Fiction"}):
    print(book["title"])

print("\nBooks published before 1950:")
for book in books_collection.find({"published_year": {"$lt": 1950}}):
    print(book["title"])

# 3. Update a document
books_collection.update_one(
    {"title": "To Kill a Mockingbird"},
    {"$set": {"in_stock": True}}
)

# 4. Delete documents
books_collection.delete_many({"genre": "Classic"})

```

## Explanation

1. **Setup:** Connects to MongoDB, gets the database 'my\_library' and collection 'books'.
2. **Insert:** Creates a list of book documents and uses `insert_many` for efficiency if you have multiple items.
3. **Query:** Demonstrates using filters to find books by genre and published year. The `$lt` is a query operator for "less than."
4. **Update:** Uses `update_one` to find a specific book and uses the `$set` operator to modify its `in_stock` status.
5. **Delete:** `delete_many` removes all documents where the genre is "Classic."

## Example 1: Data Analysis

```

import pymongo

# ... (assume we have a MongoDB connection)

# Find the average length of movies by genre
pipeline = [
    {"$group": {

```



```

        "_id": "$genre",
        "average_length": {"$avg": "$runtime_minutes" }
    }}
]
results = movies.aggregate(pipeline)

for result in results:
    print(f"{result['_id']} average length: {result['average_length']:.1f} minutes")

```

## Example 2: User Preferences

```

import pymongo

# ... (assume we have a MongoDB connection)

# Add a new movie to a user's "favorites" list
user_id = 12345
movie_id = "6258b2f053c58b732a5c06c6" # Example ObjectID

movies.update_one(
    {"_id": user_id },
    {"$addToSet": {"favorites": movie_id }}
)

```

## Example 3: Inventory Management

```

import pymongo

# ... (assume we have a MongoDB connection)

# Mark all low-stock products as "order_soon"
threshold = 10
movies.update_many(
    {"stock_quantity": {"$lt": threshold}},
    {"$set": {"status": "order_soon"}}
)

```

## Example 4: Text Search

```

import pymongo

# ... (assume we have a MongoDB connection and a text index on 'title')

search_term = "adventure"
results = movies.find({"$text": {"$search": search_term}})

for movie in results:

```

```
print(movie["title"])
```

### Example 5: Geospatial Queries (Requires Geospatial Index)

```
import pymongo

# ... (assume we have a MongoDB connection)

# Find theaters within 5 kilometers of a location
my_location = [-73.9667, 40.78] # Example: New York City
max_distance = 5000 # In meters

result = theaters.find(
    {
        "location": {
            "$near": {
                "$geometry": {"type": "Point", "coordinates": my_location},
                "$maxDistance": max_distance
            }
        }
    }
)

for theater in result:
    print(theater['name'])
```

## 4. Working with Data in MongoDB

### Setting Up MongoDB with Python

First, ensure you have the `pymongo` library installed. You can install it using pip if it's not already installed:

```
pip install pymongo
```

### 1. Data Types and Schemas

In MongoDB, data is stored in documents, which are BSON (Binary JSON) formatted. Common data types include:

- String
- Integer
- Boolean

- Date
- Array
- ObjectID (unique identifier for documents)
- etc.

MongoDB is schema-less, which means each document in a collection can have a different set of fields. However, for consistency and query efficiency, it's common to structure documents similarly within a collection.

## 2. Inserting Data

Let's insert some data into a collection called `bytes_of_intelligence`.

```
from pymongo import MongoClient

# Connection to MongoDB
client = MongoClient('localhost', 27017)
db = client['intelligence_db']
collection = db['bytes_of_intelligence']

# Insert documents
documents = [
    {"name": "Mejbah", "field": "Data Science", "technologies": ["Python", "MongoDB"], "experience_years": 2},
    {"name": "Ahammad", "field": "Machine Learning", "technologies": ["Python", "TensorFlow"], "experience_years": 3}
]
result = collection.insert_many(documents)
print(f"Inserted document ids: {result.inserted_ids}")
```

## 3. Querying Data: Basic and Advanced Queries

Now, let's query the data we just inserted, using both basic and advanced queries.

**Basic Query:**

```
# Querying a single document
mejbah = collection.find_one({"name": "Mejbah"})
print(mejbah)

# Querying multiple documents using a basic filter
developers = collection.find({"experience_years": {"$gt": 2}})
for dev in developers:
    print(dev)
```

**Advanced Query:**

```
# Using advanced query features like $in and $or
query = {
    "$or": [
        {"name": "Mejbah"},
        {"technologies": {"$in": ["TensorFlow"]}}
    ]
}
advanced_developers = collection.find(query)
for dev in advanced_developers:
    print(dev)
```

## 4. Updating and Deleting Data

Lastly, let's see how to update and delete data in the collection.

### Updating Data:

```
# Update a single document
collection.update_one(
    {"name": "Ahammad"},
    {"$set": {"experience_years": 4}}
)
print("Updated Ahammad's experience years.")

# Update multiple documents
collection.update_many(
    {"experience_years": {"$lt": 5}},
    {"$inc": {"experience_years": 1}}
)
print("Increased experience years for all developers with less than 5 years.")
```

### Deleting Data:

```
# Deleting a single document
collection.delete_one({"name": "Mejbah"})
print("Deleted Mejbah's document.")

# Deleting multiple documents
collection.delete_many({"experience_years": {"$lt": 5}})
print("Deleted all developers with less than 5 years of experience.")
```

## 5. Indexes and Performance

- Understanding Indexes

Indexes in MongoDB support the efficient execution of queries. Without indexes, MongoDB must perform a full collection scan, which can be slow if your collection has a lot of documents. Indexes store a portion of the collection's data in an easy-to-traverse form that the database can use to avoid scanning every document.

- **Creating and Managing Indexes**

### Creating an Index:

You can create indexes on a collection to improve the performance of frequent queries. The following Python code demonstrates how to create an index using `pymongo`:

```
from pymongo import MongoClient

# Connection to the MongoDB server
client = MongoClient('localhost', 27017)

# Select the database and collection
db = client['intelligence']
collection = db['bytes']

# Creating a simple single field index on the 'topic' field
collection.create_index([('topic', 1)]) # 1 for ascending order

# Creating a compound index on 'topic' and 'date'
collection.create_index([('topic', 1), ('date', -1)]) # -1 for descending order
```

### Managing Indexes:

You can list and drop indexes as needed:

```
# List all indexes on a collection
indexes = collection.list_indexes()
for index in indexes:
    print(index)

# Dropping an index
collection.drop_index('topic_1_date_-1') # Index name format is generally 'field_order'
```

- **Performance Tuning Tips**

#### 1. Use Appropriate Indexes:

- Analyze your application's query patterns and index fields that are most frequently queried.
- Remember that indexes take up extra disk space and can slow down write operations because they also need to be updated. It's a balance!

#### 2. Monitor Performance:

- Utilize tools like MongoDB's Atlas platform or the `mongostat` and `mongotop` utilities to monitor the database's performance and understand where bottlenecks might be occurring.

### 3. Index Management:

- Avoid having too many indexes. Regularly review and remove unused or less important indexes.
- Use partial indexes when you only need to index a subset of the data that meets certain criteria, which can save space and improve write performance.

### 4. Optimize Query Patterns:

- Aim to write queries that can take full advantage of the indexes.
- Use the `.explain()` method in MongoDB to analyze the efficiency of your queries.

### 5. Hardware Considerations:

- Ensure that your working set fits into RAM. MongoDB performs best when the frequently accessed data fits into memory.
- Use SSDs for better I/O performance, which is crucial for write-heavy applications.

Here's how you might use the `.explain()` method to check a query's performance:

```
# Analyzing a query's execution plan
query_performance = collection.find({"topic": "machine learning"}).explain()
print(query_performance)
```

## 6. Data Aggregation

The MongoDB Aggregation Framework is a powerful set of tools for transforming and combining data in MongoDB. It operates through a pipeline model, where data passes through multiple stages, each transforming the data in some way. It's especially useful for complex queries and reporting.

### Concepts and Usage

The aggregation pipeline is a framework in MongoDB designed to perform data transformation and summary. The data is processed as it passes through a sequence of stages, and each stage transforms the data in a specific way (e.g., filtering, grouping, sorting).

### Pipeline Stages and Common Operators

Here are a few common stages in the aggregation pipeline:

- `$match` : Filters the documents to pass only those that match the given condition into the next stage.
- `$group` : Groups documents by specified expressions.

- `$sort` : Sorts documents.
- `$project` : Passes along the documents with the requested fields to the next stage, potentially adding new fields or removing existing fields.
- `$sum` : Calculates the sum of specified values from all documents in the collection.
- `$avg` : Calculates the average of specified values from all documents.

## Practical Examples of Data Aggregation

**Setup:** Assuming we have a collection `bytes_of_intelligence` with documents structured like:

```
{
  "name": "Mejbah",
  "field": "Data Science",
  "technologies": ["Python", "MongoDB"],
  "experience_years": 5,
  "projects": 10
},
{
  "name": "Ahammad",
  "field": "Machine Learning",
  "technologies": ["Python", "TensorFlow"],
  "experience_years": 3,
  "projects": 5
}
```

### Python Code for Aggregation:

```
from pymongo import MongoClient

# Connecting to MongoDB
client = MongoClient('localhost', 27017)
db = client['intelligence_db']
collection = db['bytes_of_intelligence']

# Aggregation Pipeline
pipeline = [
    {"$match": {"technologies": "Python"}}, # Filter: Select docs with Python in technologies
    {"$group": {
        "_id": "$field", # Group by field
        "average_experience": {"$avg": "$experience_years"}, # Average experience in each field
        "total_projects": {"$sum": "$projects"} # Total projects in each field
    }},
    {"$sort": {"average_experience": -1}} # Sort by average experience descending
]

results = collection.aggregate(pipeline)
```

```
for result in results:
    print(result)
```

### Explanation:

1. **\$match** : This stage filters documents to include only those where 'technologies' array contains 'Python'.
2. **\$group** : Documents are grouped by the 'field', calculating the average years of experience and total projects for each field.
3. **\$sort** : The resulting groups are sorted based on the average experience in descending order.

## 8. Scaling and Management

As databases grow in size and demand, scaling and management become crucial for maintaining performance and availability. MongoDB provides two main mechanisms for scaling: replication and sharding. Along with these, effective monitoring and maintenance practices are vital.

### Replication: Concepts and Configuration

#### Concepts:

- **Replication** in MongoDB is the process of synchronizing data across multiple servers.
- A **replica set** is a group of MongoDB servers that maintain the same data set, providing redundancy and high availability.
- **Primary** node receives all write operations. **Secondary** nodes replicate the primary's data set and can serve read operations to increase read capacity.

**Configuration Example:** Configuring a replica set typically involves setting up multiple instances of MongoDB, then configuring them to be aware of each other. Below is a conceptual setup using Python, but keep in mind actual deployment would require MongoDB server configurations and might not involve Python code for the setup phase.

```
# This is a pseudo-code/hypothetical example for understanding concepts
from pymongo import MongoClient

# Connecting to a replica set
# The MongoClient would connect using a connection string that specifies multiple servers.
client = MongoClient('mongodb://db1.example.net:27017,db2.example.net:27017,db3.example.net:27017')

db = client['intelligence_db']
collection = db['bytes_of_intelligence']
```



```
# Inserting a document into a replica set
collection.insert_one({"name": "Mejbah", "field": "Data Science"})
```

### Notes:

- The connection string includes all members of the replica set.
- The `?replicaSet=myReplicaSet` argument specifies the name of the replica set.

## Sharding: Strategies and Implementation

### Concepts:

- **Sharding** distributes data across multiple machines to support deployments with very large data sets and high throughput operations.
- A **shard** is a single MongoDB instance that holds a portion of the dataset.
- **Shard keys** are chosen based on the document fields that determine the distribution of the collection's documents among shards.

**Implementation Example:** Setting up sharding is complex and generally handled at the database and infrastructure level. However, the concept and strategy can be outlined:

```
# Hypothetical Python code to demonstrate concept

# Assume 'shard_key' is set on a field that evenly distributes the data
# Example: Setting a shard key on 'field' if data is evenly distributed across different field
db.adminCommand({
    'shardCollection': 'intelligence_db.bytes_of_intelligence',
    'key': {'field': 1} # This would be done in the MongoDB shell, not via Python
})
```

## Monitoring and Maintenance

### Concepts:

- Monitoring involves tracking the performance, health, and availability of MongoDB instances.
- Maintenance includes tasks such as backing up data, pruning old entries, and optimizing indexes.

**Python Example:** For monitoring, you might use MongoDB's own or third-party tools, but you can access some statistics via Python:

```
# Fetching database stats
db_stats = db.command("dbstats")
print(db_stats)
```

```
# Fetching collection stats
collection_stats = db.command("collstats", "bytes_of_intelligence")
print(collection_stats)
```